


Programing in Nuprl

Vincent Rahli



May 30, 2017

Table of contents

PRL Group

Goals

Nuprl

EventML

Computational equivalence

Conclusion

PRL Group

PRL group

Robert L. Constable



Mark Bickford



Richard Eaton



Vincent Rahli



Abhishek Anand



System group

Robbert van Renesse



Nicolas Schiper



Ken Birman



Goals

Long term goal: Develop provably correct code.

Current Goals:

- ▶ Ease the programming task (EventML).
- ▶ Domain specific programming (Logic of Events/EventML).
- ▶ Generate efficient code (Constructive domain theory).

Work done as part of the CRASH project
(**Correct-by-Construction Attack-Tolerant Systems**)
funded by DARPA (Defense Advanced Research Projects Agency).

Table of Contents

PRL Group

Goals

Nuprl

EventML

Computational equivalence

Conclusion

A **constructive type theory**: CTT13 an evolution of CTT84 closely related to ITT82 [CAB⁺86, Kre02, ABC⁺06].

System **à la Curry** as opposed to Coq's system *à la* Church.

Untyped, deterministic, lazy, applied λ -calculus with:
natural numbers, pairs, injections, fix operator, \perp ,
call-by-value operator,

Nuprl

Computation System

2 meta-relations defined on top of the evaluation function [How96]:

- ▶ approximation \preceq
- ▶ computation equivalence \sim (a congruence).
$$a \sim b \triangleq a \preceq b \wedge b \preceq a.$$

Nuprl

Computation System

2 meta-relations defined on top of the evaluation function [How96]:

- ▶ approximation \preceq
- ▶ computation equivalence \sim (a congruence).
$$a \sim b \triangleq a \preceq b \wedge b \preceq a.$$

```
Inductive approxn: nat -> term -> term -> Prop :=
| sql0 : forall t1 t2, approxn 0 t1 t2
| sqlS : forall n t1 t2,
  (forall op terms1 terms2,
    computes_to t1 (CanonicalTerm op terms1)
  -> computes_to t2 (CanonicalTerm op terms2)
  -> forall a b,
    In (a,b) (combine terms1 terms2)
    -> approxn n a b)
  -> approxn (S n) t1 t2.
```

Definition approx t1 t2 := forall n, approxn n t1 t2.

For all terms t , $\perp \preceq t$.

$$\langle \perp, 1 \rangle \preceq \langle 2, 1 \rangle$$

$$\text{halts}(t) \triangleq 0 \preceq (\text{let } x := t \text{ in } 0)$$

$$\perp \sim \text{fix}(\lambda x. x).$$

Type system built on top of the untyped computation system.

A type is a **partial equivalence relation** on λ -terms [All87a, All87b].

Computational semantics: applied λ -terms provide **evidence** for the truth of propositions.

A sequent $H \vdash C \llbracket_{\text{ext}} t \rrbracket$ means that C has computational evidence (extract) t in context H .

Distributed.

Runs in the cloud.

Structured editor.

Shared library.

Tactic language: Classic ML.

Equality: $a = b \in T$

members: Ax .

Dependent function: $a:A \rightarrow B[a]$

members: f such that $\forall a \in A, f(a) \in B[a]$ 

(Extensional function equality.)

Dependent product: $a:A \times B[a]$

members: $\langle a, b \rangle$ 

Disjoint union: $A+B$

members: $\text{inl}(a), \text{inr}(b)$

Universe: \mathbb{U}_i

A hierarchy of universes to avoid Girard's paradox

Subtype: $A \sqsubseteq B$

Quotient: $T // E$

Intersection: $\cap_{a:A}. B[a]$

★Image: $\text{Img}(T, f)$

Subset: $\{a : A \mid B[a]\} \triangleq \text{Img}(a:A \times B[a], \pi_1)$

Union: $\cup_{a:A}. B[a] \triangleq \text{Img}(a:A \times B[a], \pi_2)$

Recursive type: $\text{rec}(F)$

where F is a monotone function on types [Men88].

Constructive domain theory:

Domain: Base

closed terms of the computation system quotiented by \sim

★**Approximation:** $a \preceq b$

members: Λx

Computational equivalence: $a \sim b$

members: Λx

★**Partial types:** \overline{T}

contains all members of T as well as all divergent terms

$$\text{True} \triangleq 0 \preceq 0$$

$$\text{Void} \triangleq \text{False} \triangleq 0 \preceq 1$$

$$\text{Top} \triangleq \bigcap a:\text{Void}.\text{Void}$$

$(\text{Type}, \sqsubseteq, \cap, \cup, \text{Top}, \text{Void})$ is a complete bounded lattice.

Nuprl

Timeline

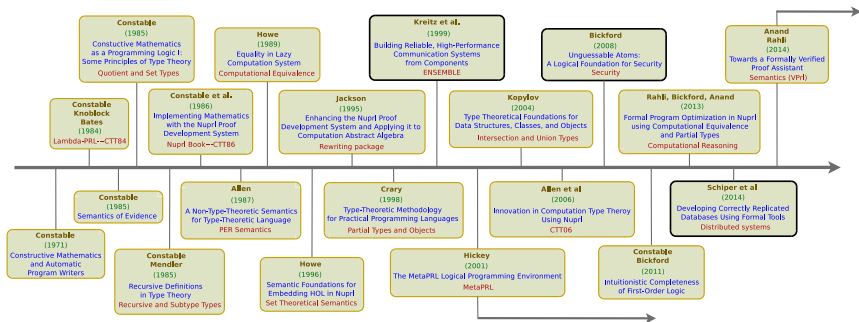


Table of Contents

PRL Group

Goals

Nuprl

EventML

Computational equivalence

Conclusion

EventML

➤ **Formal Specification, Verification, and Implementation of Fault-Tolerant Systems.**

➤ **We built a developing tool.**

ML-like: an extension of ML.

Domain-specific: asynchronous distributed programs.

Relies on a logical framework: interacts with Nuprl.

EventML

Semantics

An EventML program corresponds to both:

- ▶ a **Logic of Events** (LoE) formula
- ▶ a **General Process Model** (GPM) program.

LoE formulas are used to reason about processes.

LoE formulas are implementable by GPM programs.

EventML

Syntax

Similar to existing process calculi.

```
(* -- Leader -- *)

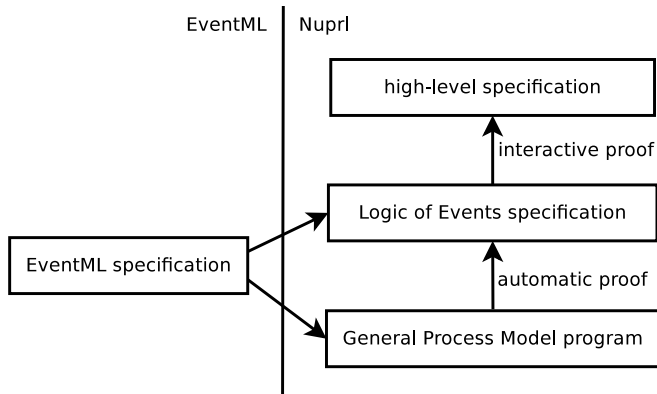
class Leader =
  (HigherLeader >>= LeaderStart)
  || (Once(start'base) >>= (\_.SpawnFirstScout))
  || ((LeaderPropose || LeaderAdopted) >>= Commander)
  || (LeaderPreempted >>= Scout)
  || ((\_.\loc.{pong'send loc ()}) o ping'base) ;;

(* ----- MAIN ----- *)

main Leader @ ldrs || Acceptor @ accpts
```

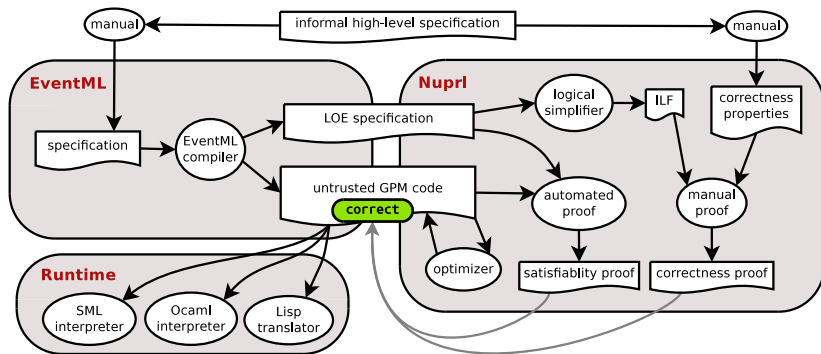
EventML

Workflow



EventML

Workflow



EventML

Verification

After spending several years on developing LoE, GPM, and EventML...

it took us:

- ▶ 2 days to prove the safety properties of 2/3-consensus,
- ▶ 3 weeks for Paxos Synod,
- ▶ 1 additional week to prove full Paxos (Synod + learners)

(Paxos is a widely used protocol to achieve software based replication which has many industrial implementations: Chubby, Megastore,)

➤ Used in ShadowDB (implemented by Nicolas Schiper).

➤ **How efficient is our generated code?**

Table of Contents

PRL Group

Goals

Nuprl

EventML

Computational equivalence

Conclusion

Computational equivalence

Motivation:

Formal program optimization in an untyped setting [RBA13].

- More general
- More efficient

Computational equivalence

A simple example:

`let $x, y = \perp$ in $x \sim \perp$?`

Computational equivalence

A simple example:

`let $x, y = \perp$ in $x \sim \perp$?`

They have the same observable behavior.

How can we prove this equivalence?

Computational equivalence

A simple example:

$\text{let } x, y = \perp \text{ in } x \sim \perp?$

They have the same observable behavior.

How can we prove this equivalence?

We have to prove:

$\text{let } x, y = \perp \text{ in } x \preceq \perp$

$\perp \preceq \text{let } x, y = \perp \text{ in } x$

Computational equivalence

$\perp \preceq \text{let } x, y = \perp \text{ in } x$ is trivial.

How about:

$\text{let } x, y = \perp \text{ in } x \preceq \perp$

By definition of \preceq we can assume:

$\text{halts}(\text{let } x, y = \perp \text{ in } x)$

We added a rule that says:

if $\text{halts}(\text{let } x, y = t \text{ in } F)$ then $t \sim \langle \pi_1(t), \pi_2(t) \rangle$

(And similarly for all destructors.)

Computational equivalence

⇒ **We added rules to reason about the computation system**

Computational equivalence

A more complicated example:

$$\forall t : \text{Top}. \text{map}(f, \text{map}(g, t)) \sim \text{map}(f \circ g, t)?$$

$$\begin{aligned} & \text{map}(f, t) \\ &= \text{fix} \left(\lambda R. \lambda t. \text{ispair} \left(\begin{array}{l} t, \\ \text{let } x, y = t \text{ in } (f \ x) \bullet R \ y, \\ \text{isaxiom}(t, \text{nil}, \perp) \end{array} \right) \right) t \end{aligned}$$

a list: $\langle 1, \langle 2, \langle 3, \text{Ax} \rangle \rangle \rangle$

Computational equivalence

⇒ We added the following least upper bound property [Cra98]

$$\begin{array}{l} H \vdash G[\text{fix}(f)/x] \preceq t \\ \text{BY } [\text{least-upper-bound}] \\ H, n : \mathbb{N} \vdash G[f^n(\perp)/x] \preceq t \end{array}$$

We prove

$$\text{map}(f \circ g, t) \preceq \text{map}(f, \text{map}(g, t))$$

using [least-upper-bound] and then by induction on n .

Computational equivalence

In the induction case, we end up with:

$$\text{ispair} \left(\begin{array}{l} t, \\ \text{let } x, y = t \text{ in } (f\ x) \bullet R\ y, \\ \text{isaxiom}(t, \text{nil}, \perp) \end{array} \right) \preceq X$$

➤ **We added the following rule:**

$$H \vdash C \text{ [ext ispair}(t, a, b)[x \backslash Ax]]$$

BY [ispairCases]

$$H \vdash \text{halts}(t)$$
$$H \vdash t \in \text{Base}$$
$$H, x : t \sim \langle \pi_1(t), \pi_2(t) \rangle \vdash C \text{ [ext } a]$$
$$H, x : (\forall [u, v : \text{Base}]. \text{ispair}(z, u, v) \sim v)[z \backslash t] \vdash C \text{ [ext } b]$$

Computational equivalence

Process type:

$$\text{corec}(\lambda P. A \rightarrow P \times \text{Bag}(B))$$

where

$$\text{corec}(G) = \bigcap_{n:\mathbb{N}} \text{fix} \left(\begin{array}{l} \lambda P. \lambda n. \text{if } n =_{\mathbb{Z}} 0 \text{ then Top} \\ \text{else } G(P(n-1)) \end{array} \right) n$$

$$P = \text{buffer}((\lambda n. \lambda \text{buf}. \{n + \text{buf}\}) \circ \text{base}(\lambda m. \{m\}), \{0\})$$

\Downarrow

$$P' = \text{fix}(\lambda F. \lambda s. \lambda m. \text{let } x ::= m + s \text{ in } \langle F\ x, \{x\} \rangle) 0$$

Computational equivalence

⇒ P vs. P' :

- ▶ 100/200 computation steps for P
- ▶ less than 10 computation steps for P'

⇒ ShadowDB:

- ▶ non-optimized code: 127 milliseconds
- ▶ optimized code: 60 milliseconds
- ▶ Lisp code: 5 milliseconds
- ▶ reference implementation: 1 millisecond

Current and future work

➤ Performance

- ▶ Identify more optimizations.
- ▶ Prove that our optimizations improve the runtime.
- ▶ Translators and compiler?

➤ ShadowDB

- ▶ Deal with Byzantine faults.
- ▶ Synthesize more of ShadowDB.

➤ Nuprl

- ▶ Prove that our new types and rules are valid.

References |



Stuart F. Allen, Mark Bickford, Robert L. Constable, Richard Eaton, Christoph Kreitz, Lori Lorigo, and Evan Moran.
Innovations in computational type theory using Nuprl.
J. Applied Logic, 4(4):428–469, 2006.
<http://www.nuprl.org/>.



Stuart F. Allen.
A non-type-theoretic definition of Martin-Löf's types.
In *LICS*, pages 215–221. IEEE Computer Society, 1987.



Stuart F. Allen.
A Non-Type-Theoretic Semantics for Type-Theoretic Language.
PhD thesis, Cornell University, 1987.



R.L. Constable, S.F. Allen, H.M. Bromley, W.R. Cleaveland, J.F. Cremer, R.W. Harper, D.J. Howe, T.B. Knoblock, N.P. Mendler, P. Panangaden, J.T. Sasaki, and S.F. Smith.
Implementing mathematics with the Nuprl proof development system.
Prentice-Hall, Inc., Upper Saddle River, NJ, USA, 1986.



Karl Crary.
Type-Theoretic Methodology for Practical Programming Languages.
PhD thesis, Cornell University, Ithaca, NY, August 1998.



Douglas J. Howe.
Proving congruence of bisimulation in functional programming languages.
Inf. Comput., 124(2):103–112, 1996.

References II



Christoph Kreitz.

The Nuprl Proof Development System, Version 5, Reference Manual and User's Guide.

Cornell University, Ithaca, NY, 2002.

www.nuprl.org/html/02cucs-NuprlManual.pdf.



Paul F. Mendler.

Inductive Definition in Type Theory.

PhD thesis, Cornell University, Ithaca, NY, 1988.



Vincent Rahli, Mark Bickford, and Abhishek Anand.

Formal program optimization in Nuprl using computational equivalence and partial types.

In *ITP'13*, volume 7998 of *LNCS*, pages 261–278. Springer, 2013.